

# 1 Benchmarking Program

This section describes the usage of the benchmarking program and its source code. XXX what does the benchmark program do? The structure of the source file is listed below.

```
1a  <bench.cpp 1a>≡
    #include <cstdio>
    #include <cstdlib>
    #include <cstring>
    #include <ctime>
    #include <stdint>
    #include <sycl/sycl.hpp>
    #include <omp.h>
    <Additional headers 12a>
    <Macro definitions 2a>
    <Type definitions 1b>
    <Function prototypes 2d>
    <Global variables 2b>
    <Function definitions 3d>
    int main(int argc, const char **argv)
    {
        <Process arguments 2c>
        <Possibly initialize dump 5c>
        <Time the benchmark 4b>
        <Print the output 5a>
        <Possibly finalize dump 5d>
        return 0;
    }
```

Defines:

`argc`, used in chunk 2c.  
`argv`, used in chunks 2–4.  
`main`, never used.

The program supports the following hash algorithms and methods to generate the hashes.

```
1b  <Type definitions 1b>≡ (1a)
    enum algorithm {SHA224, SHA256, BLAKE3};
    enum method {SERIAL, OPENMP, SYCL_CPU, SYCL_GPU};
```

Uses `algorithm 2b` and `method 2b`.

The following `printf` template is used for printing output to the standard output.

2a  $\langle$ Macro definitions 2a $\rangle \equiv$  (1a) 3b $\triangleright$

```
#define OUTPUT_TEMPLATE "hashes_per_block =\t%u\t" \
    "num_blocks =\t%u\t" \
    "algorithm =\t%s\t" \
    "runner =\t%s\t" \
    "elapsed (s) =\t%f\n"
```

Uses `algorithm` 2b, `elapsed` 4b, `hashes_per_block` 2b, and `num_blocks` 2b.

## 1.1 Argument Processing

This subsection details the code for processing `argv`. The arguments listed above are stored in the following global variables, respectively. Algorithms and methods are stored using enumerations defined earlier.

2b  $\langle$ Global variables 2b $\rangle \equiv$  (1a) 3a $\triangleright$

```
static unsigned hashes_per_block;
static unsigned num_blocks;
static enum algorithm algorithm;
static enum method method;
```

Defines:

`algorithm`, used in chunks 1–3, 5–8, and 11–13.  
`hashes_per_block`, used in chunks 2a, 3c, 5a, 7, 8, and 11a.  
`method`, used in chunks 1b, 3c, 5a, and 7a.  
`num_blocks`, used in chunks 2a, 3c, 5a, 8, and 11a.

If the incorrect number of arguments were given, then we print an error message and exit. Currently, we are assuming that the user will have access to the documentation and will be able to find the usage information.

2c  $\langle$ Process arguments 2c $\rangle \equiv$  (1a) 3c $\triangleright$

```
if (argc != 5) {
    fprintf(stderr, "%s: incorrect number of arguments.\n", argv[0]);
    return 1;
}
```

Uses `argc` 1a and `argv` 1a.

Processing each argument takes several function calls. Also, if there is an error processing the arguments, we want to exit with an error. The below functions will handle these tasks. They both take `argv` and the argument within `argv` to parse. The `parse_enumerator` function also takes an array of valid enumerators and the number of valid enumerators; the parsed value is the index of the enumerator within the array. Both functions return the parsed value.

2d  $\langle$ Function prototypes 2d $\rangle \equiv$  (1a) 6b $\triangleright$

```
static unsigned parse_unsigned(const char **argv, int arg_num);
static int parse_enumerator(const char **argv, int arg_num,
    const char **enumerators, const unsigned num_enumerators);
```

Uses `argv` 1a, `parse_enumerator` 4a, and `parse_unsigned` 3d.

The algorithm and method enumerations both need an array of enumerators for `parse_enumerator`. The order of strings within the array **must** match the order of enumerators within the enumeration definition.

- 3a  $\langle \textit{Global variables 2b} \rangle + \equiv$  (1a)  $\triangleleft 2b \ 5b \triangleright$   
`static const char *algorithms[] = {"sha224", "sha256", "blake3"};`  
`static const char *methods[] = {"serial", "openmp", "sycl-cpu", "sycl-gpu"};`
- It will be useful to have a macro to take the length of an array.
- 3b  $\langle \textit{Macro definitions 2a} \rangle + \equiv$  (1a)  $\triangleleft 2a$   
`#define LENGTH(arr) (sizeof(arr) / sizeof((arr)[0]))`
- C++ wants additional casts for enumerators.
- 3c  $\langle \textit{Process arguments 2c} \rangle + \equiv$  (1a)  $\triangleleft 2c$   
`hashes_per_block = parse_unsigned(argv, 1);`  
`num_blocks = parse_unsigned(argv, 2);`  
`algorithm = (enum algorithm)`  
`parse_enumerator(argv, 3, algorithms, LENGTH(algorithms));`  
`method = (enum method)parse_enumerator(argv, 4, methods, LENGTH(algorithms));`

Uses `algorithm 2b`, `argv 1a`, `hashes_per_block 2b`, `method 2b`, `num_blocks 2b`, `parse_enumerator 4a`, and `parse_unsigned 3d`.

Parsing integers is simple using C's `sscanf`. If it does not match any inputs items, which is the only case for failure here, it returns EOF. It cannot fail before the first match, because there is only one item to match.

- 3d  $\langle \textit{Function definitions 3d} \rangle \equiv$  (1a)  $4a \triangleright$   
`static unsigned parse_unsigned(const char **argv, int arg_num)`  
`{`  
 `int rc;`  
 `unsigned result;`  
 `rc = sscanf(argv[arg_num], "%u", &result);`  
 `if (rc == EOF) {`  
 `fprintf(stderr, "%s: could not parse \"%s\" as an unsigned "`  
 `"integer\n", argv[0], argv[arg_num]);`  
 `exit(1);`  
 `}`  
 `return result;`  
`}`

Defines:

`parse_unsigned`, used in chunks 2d and 3c.

Uses `argv 1a`.

Enumerators are parsed using a linear scan and C's `strcmp`. There are ways to make this more efficient, but it probably does not matter. Printing the list of enumerators as part of the error message is non-trivial, and possibly unnecessary.

```
4a  <Function definitions 3d>+≡ (1a) <3d 5e>
    static int parse_enumerator(const char **argv, int arg_num,
                                const char **enumerators, const unsigned num_enumerators)
    {
        unsigned i;
        for (i = 0; i < num_enumerators; i++)
            if (strcmp(argv[arg_num], enumerators[i]) == 0)
                return i;
        fprintf(stderr, "%s: could not match \"%s\" to ",
                argv[0], argv[arg_num]);
        for (i = 0; i < num_enumerators - 1; i++)
            fprintf(stderr, "\"%s", enumerators[i]);
        fprintf(stderr, "or \"%s\".\n", enumerators[i]);
        exit(1);
        return 0;
    }
```

Defines:

`parse_enumerator`, used in chunks 2d and 3c.

Uses `argv` 1a.

## 1.2 Timing the Benchmark

The POSIX clock interface is used to get the time before and after the hashes are generated. Since only a relative time is required, `CLOCK_MONOTONIC` is sufficient. The elapsed time is stored in the variable `elapsed`.

```
4b  <Time the benchmark 4b>≡ (1a)
    double elapsed;
    struct timespec start, end;
    clock_gettime(CLOCK_MONOTONIC, &start);
    <Run the benchmark 7a>
    clock_gettime(CLOCK_MONOTONIC, &end);
    elapsed = (double)end.tv_sec - (double)start.tv_sec;
    elapsed += ((double)end.tv_nsec - (double)start.tv_nsec) / 1e12L;
```

Defines:

`elapsed`, used in chunks 2a and 5a.

### 1.3 Printing the Output

All of the variables that are needed for output have been defined. The string names for the algorithm and method are printed instead of their enumerator's number.

5a     $\langle$ *Print the output 5a* $\rangle \equiv$  (1a)  
       printf(OUTPUT\_TEMPLATE, hashes\_per\_block, num\_blocks,  
               algorithms[algorithm], methods[method], elapsed);  
 Uses algorithm 2b, elapsed 4b, hashes\_per\_block 2b, method 2b, and num\_blocks 2b.

### 1.4 Dumping Hashes

If the DUMP preprocessor macro is defined, the program should dump its hashes to a file. This is for validating that the generated hashes are correct.

See the check-dumps.sh script to compare the hashes side-by-side.

5b     $\langle$ *Global variables 2b* $\rangle + \equiv$  (1a)  $\langle$ 3a 7d $\rangle$   
       #ifdef DUMP  
       static FILE \*dump\_stream;  
       #endif

5c     $\langle$ *Possibly initialize dump 5c* $\rangle \equiv$  (1a)  
       #ifdef DUMP  
       dump\_stream = fopen("bench-hashes.dat", "w");  
       #endif

5d     $\langle$ *Possibly finalize dump 5d* $\rangle \equiv$  (1a)  
       #ifdef DUMP  
       fclose(dump\_stream);  
       #endif

All of the generators will call `dump`. When dumping is enabled, it will be an actual function. Otherwise, it will be a macro that expands to nothing.

5e     $\langle$ *Function definitions 3d* $\rangle + \equiv$  (1a)  $\langle$ 4a 6c $\rangle$   
       #ifdef DUMP  
       static void dump(unsigned char \*buffer, size\_t num\_hashes, size\_t hash\_size)  
       {  
            $\langle$ *Dump hashes in buffer 6a* $\rangle$   
       }  
       #else  
       #define dump(x, y, z) /\* dump \*/  
       #endif

Defines:

dump, used in chunks 8 and 11a.

Each hash is printed in hexadecimal, on its own line. TODO: it might be useful to print the index along with the hash.

6a  $\langle \text{Dump hashes in buffer 6a} \rangle \equiv$  (5e)

```

    unsigned i, j;
    for (i = 0; i < num_hashes; i++) {
        for (j = 0; j < hash_size; j++) {
            fprintf(dump_stream, "%02x", buffer[i * hash_size + j]);
        }
        fprintf(dump_stream, "\n");
    }

```

## 1.5 Supporting Different Hash Algorithms

The goal is to support several different hash algorithms and compare their performance. The following function will be used to dispatch the correct hash function according the algorithm, and ensure the result is written to `buf[slot]`. The algorithm must be passed explicitly, otherwise the SYCL compiler will complain.

6b  $\langle \text{Function prototypes 2d} \rangle + \equiv$  (1a)  $\triangleleft 2d$

```

    static void run_hash(enum algorithm algorithm, uint64_t input,
        unsigned char *buf, unsigned slot);

```

Uses `algorithm 2b` and `run_hash 6c`.

The downside of this approach is that the algorithm is checked every time a hash is generated. This could be avoided by calling a function that is set to the appropriate hash function, but SYCL does not support calling function pointers in its kernels.

6c  $\langle \text{Function definitions 3d} \rangle + \equiv$  (1a)  $\triangleleft 5e \triangleright 9$

```

    static void run_hash(enum algorithm algorithm, uint64_t input,
        unsigned char *buf, unsigned slot)
    {
        switch (algorithm) {
        case SHA224:
             $\langle \text{Hash input to buf[slot] with SHA-224 12b} \rangle$ 
            break;
        case SHA256:
             $\langle \text{Hash input to buf[slot] with SHA-256 12e} \rangle$ 
            break;
        case BLAKE3:
             $\langle \text{Hash input to buf[slot] with BLAKE3 13b} \rangle$ 
            break;
        }
    }

```

Defines:

`run_hash`, used in chunks 6b, 8, and 11a.

Uses `algorithm 2b`.

## 1.6 Supporting Different Running Methods

The other feature of this program is that it supports several drivers to run the has generation code. Some runners may require common, local variables.

7a  $\langle \text{Run the benchmark 7a} \rangle \equiv$  (4b)

```

     $\langle \text{Local declarations for the runner 7b} \rangle$ 
    switch (method) {
    case SERIAL:
         $\langle \text{Run benchmark in serial 8a} \rangle$ 
        break;
    case SYCL_CPU:
         $\langle \text{Run benchmark with SYCL on the CPU 11b} \rangle$ 
        break;
    case SYCL_GPU:
         $\langle \text{Run benchmark with SYCL on the GPU 11c} \rangle$ 
        break;
    case OPENMP:
         $\langle \text{Run benchmark with OpenMP 8b} \rangle$ 
        break;
    }
     $\langle \text{Delete any local declarations 7c} \rangle$ 

```

Uses method 2b.

Something that all of the runners will need is a buffer to write the hashes to. This requires knowing the size of each hash in bytes. They are stored in the `digest_size` array.

7b  $\langle \text{Local declarations for the runner 7b} \rangle \equiv$  (7a)

```

    unsigned char *output_buffer =
        new unsigned char[hashes_per_block * digest_size[algorithm]];

```

Uses algorithm 2b and hashes\_per\_block 2b.

7c  $\langle \text{Delete any local declarations 7c} \rangle \equiv$  (7a)

```

    delete[] output_buffer;

```

7d  $\langle \text{Global variables 2b} \rangle + \equiv$  (1a) <5b

```

    static const unsigned digest_size[] = {28u, 32u, 32u};

```

## 1.7 Running in Serial

Running in serial is simple.

8a  $\langle \textit{Run benchmark in serial 8a} \rangle \equiv$  (7a)

```

    for (uint64_t i = 0; i < num_blocks; i++) {
        for (uint64_t j = 0; j < hashes_per_block; j++) {
            run_hash(algorithm, i * hashes_per_block + j, output_buffer, j);
        }
        dump(output_buffer, hashes_per_block, digest_size[algorithm]);
    }

```

Uses `algorithm 2b`, `dump 5e`, `hashes_per_block 2b`, `num_blocks 2b`, and `run_hash 6c`.

## 1.8 Running with OpenMP

The same as serial, but with an OpenMP pragma.

There are probably other ways to parallelize this.

8b  $\langle \textit{Run benchmark with OpenMP 8b} \rangle \equiv$  (7a)

```

#pragma omp parallel
for (uint64_t i = 0; i < num_blocks; i++) {
#pragma omp for
    for (uint64_t j = 0; j < hashes_per_block; j++) {
        run_hash(algorithm, i * hashes_per_block + j, output_buffer, j);
    }
    dump(output_buffer, hashes_per_block, digest_size[algorithm]);
}

```

Uses `algorithm 2b`, `dump 5e`, `hashes_per_block 2b`, `num_blocks 2b`, and `run_hash 6c`.



## 1.9 Running with SYCL

The following function looks for devices that match the selector, and returns queues for them. If `use_all` is false, then only one queue is returned.

```

9  <Function definitions 3d>+≡ (1a) <6c 10>
    template<class Selector>
    static std::vector<sycl::queue> make_queues(Selector sel, bool use_all)
    {
        sycl::platform p(sel);
        std::vector<sycl::device> ds = p.get_devices();
        if (!use_all) {
            for (unsigned i = 1; i < ds.size(); i++) {
                ds.pop_back();
            }
        }
        std::vector<sycl::queue> result;
        for (sycl::device d : ds) {
            sycl::queue q(d);
            result.push_back(q);
        }
        return result;
    }

```

Defines:

`make_queues`, used in chunk 11a.

The following function creates buffers for each device in the given vector. If allocation fails, the program exits with an error.

TODO: better error message

```

10  <Function definitions 3d>+≡ (1a) <9 11a>
    static std::vector<unsigned char *> alloc_buffers(std::vector<sycl::queue> qs,
        int buffer_size)
    {
        std::vector<unsigned char *> result;
        for (sycl::queue q : qs) {
            unsigned char *b = sycl::malloc_device<unsigned char>(
                buffer_size, q);
            if (b == nullptr) {
                fprintf(stderr, "sycl::malloc_device failed when "
                    "called %u bytes were requested.\n",
                        buffer_size);
                exit(1);
            }
            result.push_back(b);
        }
        return result;
    }

```

Defines:

alloc\_buffers, used in chunk 11a.

Running with SYCL requires a few local variables and polymorphic types. A separate function is used to deal with this.

TODO: document `use_all`.

11a  $\langle \text{Function definitions 3d} \rangle \equiv$  (1a)  $\triangleleft 10$

```

template <class Selector>
static void run_sycl(Selector sel, unsigned char *host_buffer, bool use_all)
{
    std::vector<sycl::queue> qs = make_queues(sel, use_all);
    int buffer_size = (hashes_per_block * digest_size[algorithm])
        / qs.size();
    int hashes_per_device = hashes_per_block / qs.size();
    std::vector<unsigned char *> buffers = alloc_buffers(qs, buffer_size);
#pragma omp parallel if(qs.size() > 1) num_threads(qs.size())
    {
        int t = omp_get_thread_num();
        enum algorithm alg = algorithm;
        unsigned char *host_ptr = host_buffer + buffer_size * t;
        for (uint64_t i = 0, base = 0; i < num_blocks;
             i++, base += hashes_per_device) {
            sycl::event hashes_ev = qs[t].parallel_for(sycl::range<1>(hashes_per_block),
                run_hash(alg, base + idx, buffers[t], idx));
        };
        sycl::event copy_ev = qs[t].memcpy(host_ptr, buffers[t], buffer_size);
        copy_ev.wait();
        dump(host_ptr, hashes_per_device, digest_size[algorithm]);
    }

    for (unsigned i = 0; i < buffers.size(); i++) {
        sycl::free(buffers[i], qs[i]);
    }
}

```

Defines:

`run_sycl`, used in chunk 11.

Uses `algorithm` 2b, `alloc_buffers` 10, `dump` 5e, `hashes_per_block` 2b, `make_queues` 9, `num_blocks` 2b, and `run_hash` 6c.

TODO runners that set `use_all` to false.

11b  $\langle \text{Run benchmark with SYCL on the CPU 11b} \rangle \equiv$  (7a)

```
run_sycl(sycl::cpu_selector_v, output_buffer, true);
```

Uses `run_sycl` 11a.

11c  $\langle \text{Run benchmark with SYCL on the GPU 11c} \rangle \equiv$  (7a)

```
run_sycl(sycl::gpu_selector_v, output_buffer, true);
```

Uses `run_sycl` 11a.

## 1.10 SHA-224 Hash Algorithm

The SHA-224 algorithm used here comes from an implementation that I found online.<sup>1</sup> There are some optimizations that could be made with a custom implementation, because we know exactly how long the message is, etc.

12a  $\langle \text{Additional headers 12a} \rangle \equiv$  (1a) 12c  $\triangleright$   
`#include "sha224.hpp"`

Using this implementation requires only the following function calls. The hash algorithms are implemented in `switch` case bodies, so a new scope is needed to declare local variables well.

12b  $\langle \text{Hash input to buf[slot] with SHA-224 12b} \rangle \equiv$  (6c)  

```
{
    class SHA224 ctx;
    ctx.init();
    ctx.update((const unsigned char *)&input, sizeof(input));
    ctx.final(buf + slot * digest_size[algorithm]);
}
```

Uses algorithm 2b.

As it turns out SYCL\_EXTERNAL is not universal, and everything **must** be in the same translation to be portable across SYCL compilers, it seems. Although the below strategy works, it's bad practice and should be replaced.

12c  $\langle \text{Additional headers 12a} \rangle + \equiv$  (1a)  $\triangleleft$  12a 12d  $\triangleright$   
`#include "sha224.cpp"`

## 1.11 SHA-256 Hash Algorithm

I modified the SHA-224 implementation to be SHA-256. I have not completely verified that it is correct.

12d  $\langle \text{Additional headers 12a} \rangle + \equiv$  (1a)  $\triangleleft$  12c 12f  $\triangleright$   
`#include "sha256.hpp"`

12e  $\langle \text{Hash input to buf[slot] with SHA-256 12e} \rangle \equiv$  (6c)  

```
{
    class SHA256 ctx;
    ctx.init();
    ctx.update((const unsigned char *)&input, sizeof(input));
    ctx.final(buf + slot * digest_size[algorithm]);
}
```

Uses algorithm 2b.

See the SHA-224 section for an explanation.

12f  $\langle \text{Additional headers 12a} \rangle + \equiv$  (1a)  $\triangleleft$  12d 13a  $\triangleright$   
`#include "sha256.cpp"`

---

<sup>1</sup><http://www.zedwood.com/article/cpp-sha224-function>

## 1.12 BLAKE3 Hash Algorithm

The BLAKE3 implementation comes from the BLAKE3 reference implementation. It had to be modified slightly to work with C++ and SYCL.

13a  $\langle \text{Additional headers 12a} \rangle + \equiv$  (1a)  $\triangleleft 12f \ 13c \triangleright$   
`#include "blake3.h"`

The interface for BLAKE3 uses different names and arguments than SHA-2, but functionally the same otherwise.

13b  $\langle \text{Hash input to buf[slot] with BLAKE3 13b} \rangle \equiv$  (6c)  

```
{
    blake3_hasher hasher;
    blake3_hasher_init(&hasher);
    blake3_hasher_update(&hasher, &input, sizeof(input));
    blake3_hasher_finalize(&hasher,
                          buf + slot * digest_size[algorithm],
                          digest_size[algorithm]);
}
```

Uses `algorithm 2b`.

See the SHA-224 section for an explanation.

13c  $\langle \text{Additional headers 12a} \rangle + \equiv$  (1a)  $\triangleleft 13a$   
`#include "blake3.cpp"`  
`#include "blake3_dispatch.cpp"`  
`#include "blake3_portable.cpp"`

## 1.13 Support Scripts

The `run-bench.sh` script automates the collection of data using the benchmark, and outputs the results to the file `bench-results`. It generates SHA-256 and BLAKE3 hashes on the CPU and GPU with block sizes and hash counts read from the standard input. The `generate-inputs.sh` script, described later, generates suitable inputs.

It fails for the huge buffer sizes. This should be investigated, but it does not appear to affect the completion of the script or the results.

13d  $\langle \text{run-bench.sh 13d} \rangle \equiv$   

```
#!/bin/sh

while read -r x y
do
    ./bench $x $y sha256 openmp
    ./bench $x $y sha256 sycl-cpu
    ./bench $x $y sha256 sycl-gpu
    ./bench $x $y blake3 openmp
    ./bench $x $y blake3 sycl-cpu
    ./bench $x $y blake3 sycl-gpu
done | tee bench-results
```

## 1.14 Making Graphs

The following script can be used to convert the output of the benchmark script to a `grap(1)` graph, suitable to `copy` in `report.ms`. See the report source of that file to see how to use the output of this script.

This script needs some work, especially with selecting ticks.

Assumes 256-bit hashes.

```
14  <to-grap.sh 14>≡
    #!/bin/sh
    echo 'label left "Real Time" "(s)"'
    echo 'label bot "Block Size"'
    awk '-F\t' '
        function tosize(n, i, suffix, suffixes) {
            <Convert bytes to human-readable units 15a>
        }
        BEGIN {
            xmin = 1e12
        }
        /openmp/ && /sha/ {char[count] = "ci"}
        /cpu/ && /sha/ {char[count] = "sq"}
        /gpu/ && /sha/ {char[count] = "pl"}
        /openmp/ && /blake/ {char[count] = "bu"}
        /cpu/ && /blake/ {char[count] = "*D"}
        /gpu/ && /blake/ {char[count] = "mu"}
        {
            x[count] = $2
            y[count] = $10
            ymax = $10 > ymax ? $10 : ymax
            xmax = $2 > xmax ? $2 : xmax
            xmin = $2 < xmin ? $2 : xmin
            xtick[$2]++
            count++
        }
        END {
            printf "coord y 0, %d log x\n", ymax + 1
            <Select and print ticks 15b>
            printf "\"\\(sq SHA-256, CPU\" \"\\(pl SHA-256, GPU\" \" \" \" \\
                \"\\(*D BLAKE3, CPU\" \"\\(mu BLAKE3, GPU\" \" \" \" \\
                \"\\(ci SHA256, OMP\" \"\\(bu BLAKE3, OMP\" \" \" \" \\
                \"\\\" ljust at (%f,%f)\n\", xmax / 8, ymax * 0.8
            for (i = 0; i < count; i++) {
                if (!(i in char)) continue
                printf "\"\\(%s\" at (%f,%f)\n\", char[i], x[i], y[i]
            }
        }
    ,
```

TODO is there an easier way to do this with gawk?

15a  $\langle \text{Convert bytes to human-readable units 15a} \rangle \equiv$  (14)

```

    suffixes[0] = "B"
    suffixes[1] = "KiB"
    suffixes[2] = "MiB"
    suffixes[3] = "GiB"
    i = 0
    suffix = 0
    while (n > 1024) {
        n = n / 1024
        suffix++
    }
    return sprintf("%d%s", n, suffixes[suffix])

```

TODO improve this. This seems bad.

15b  $\langle \text{Select and print ticks 15b} \rangle \equiv$  (14)

```

    printf "ticks bot out at"
    len = asorti(xtick)
    comma = 0
    for (i = 0; i < len; i += 2) {
        if (xtick[i] == 0) {
            i--
            continue
        }
        printf "%s%d \"%s\"", comma ? ", " : " ", xtick[i], tosize(xtick[i] * 32)
        comma = 1
    }
    printf "\n"

```

## 1.15 Generating Inputs

The following script generates inputs for run-bench.sh. It takes two arguments: the amount of bytes to generate and the maximum block size. It assumes 256-bit hashes.

Arguments are taken as powers of two. For example, to generate 16GiB of hashes with a maximum block size of 4GiB, use the command line below. Note that  $2^{34} = 16\text{GiB}$  and  $2^{31} = 4\text{GiB}$ .

```
./generate-inputs.sh 34 31
```

```
16  <generate-inputs.sh 16>≡
    #!/bin/sh

    if [ $# -lt 2 ]
    then
        printf 'usage: ./generate-inputs.sh TOTAL_SIZE MAX_BLOCK_SIZE\n' >&2
        printf '\tArguments are powers of 2, e.g., 32 gives 2^32 = 16GiB\n' >&2
        printf '\t2^0=1 2^1=2 2^2=4 2^3=8 2^4=16 2^5=32 2^6=64 2^7=128 ' >&2
        printf '2^8=256 2^9=512\n' >&2
        exit 1
    fi

    total_bytes=$1
    max_block_bytes=$2
    digest_size_bytes=5 # 256 bits = 32 bytes, 32=2^5

    total=$((total_bytes - digest_size_bytes))
    max_block=$((max_block_bytes - digest_size_bytes))

    for b in $(seq 10 $max_block)
    do
        printf '%d\t%d\n' $((1 << b)) $((1 << (total - b)))
    done
```



## 1.16 Verifying Hashes

This script generates dumps of hashes generated on the CPU and GPU, then displays them side-by-side with `less(1)`.

```
17  <check-dumps.sh 17>≡
    #!/bin/sh
    ./bench 1024 1024 blake3 serial
    mv bench-hashes.dat bench-hashes-serial.txt
    ./bench 1024 1024 blake3 sycl-cpu
    mv bench-hashes.dat bench-hashes-cpu.txt
    ./bench 1024 1024 blake3 sycl-gpu
    mv bench-hashes.dat bench-hashes-gpu.txt
    paste bench-hashes-serial.txt bench-hashes-cpu.txt bench-hashes-gpu.txt | less
    wc -l bench-hashes-serial.txt bench-hashes-cpu.txt bench-hashes-gpu.txt
```

## 2 Index

### 2.1 Chunks

⟨Additional headers 12a⟩  
 ⟨bench.cpp 1a⟩  
 ⟨check-dumps.sh 17⟩  
 ⟨Convert bytes to human-readable units 15a⟩  
 ⟨Delete any local declarations 7c⟩  
 ⟨Dump hashes in `buffer` 6a⟩  
 ⟨Function definitions 3d⟩  
 ⟨Function prototypes 2d⟩  
 ⟨generate-inputs.sh 16⟩  
 ⟨Global variables 2b⟩  
 ⟨Hash input to `buf[slot]` with *BLAKE3* 13b⟩  
 ⟨Hash input to `buf[slot]` with *SHA-224* 12b⟩  
 ⟨Hash input to `buf[slot]` with *SHA-256* 12e⟩  
 ⟨Local declarations for the runner 7b⟩  
 ⟨Macro definitions 2a⟩  
 ⟨Possibly finalize dump 5d⟩  
 ⟨Possibly initialize dump 5c⟩  
 ⟨Print the output 5a⟩  
 ⟨Process arguments 2c⟩  
 ⟨Run benchmark in serial 8a⟩  
 ⟨Run benchmark with *OpenMP* 8b⟩  
 ⟨Run benchmark with *SYCL* on the CPU 11b⟩  
 ⟨Run benchmark with *SYCL* on the GPU 11c⟩  
 ⟨Run the benchmark 7a⟩  
 ⟨run-bench.sh 13d⟩  
 ⟨Select and print ticks 15b⟩  
 ⟨Time the benchmark 4b⟩  
 ⟨to-grap.sh 14⟩  
 ⟨Type definitions 1b⟩

### 2.2 Identifiers

algorithm: 1b, 2a, 2b, 3c, 5a, 6b, 6c, 7b, 8a, 8b, 11a, 12b, 12e, 13b  
 alloc\_buffers: 10, 11a  
 argc: 1a, 2c  
 argv: 1a, 2c, 2d, 3c, 3d, 4a  
 dump: 5e, 8a, 8b, 11a  
 elapsed: 2a, 4b, 5a  
 hashes\_per\_block: 2a, 2b, 3c, 5a, 7b, 8a, 8b, 11a  
 main: 1a  
 make\_queues: 9, 11a  
 method: 1b, 2b, 3c, 5a, 7a

num\_blocks: 2a, 2b, 3c, 5a, 8a, 8b, 11a  
parse\_enumerator: 2d, 3c, 4a  
parse\_unsigned: 2d, 3c, 3d  
run\_hash: 6b, 6c, 8a, 8b, 11a  
run\_sycl: 11a, 11b, 11c